

COMP90038

Algorithms and Complexity

Lecture 13: Priority Queues, Heaps and Heapsort
(with thanks to Harald Søndergaard)

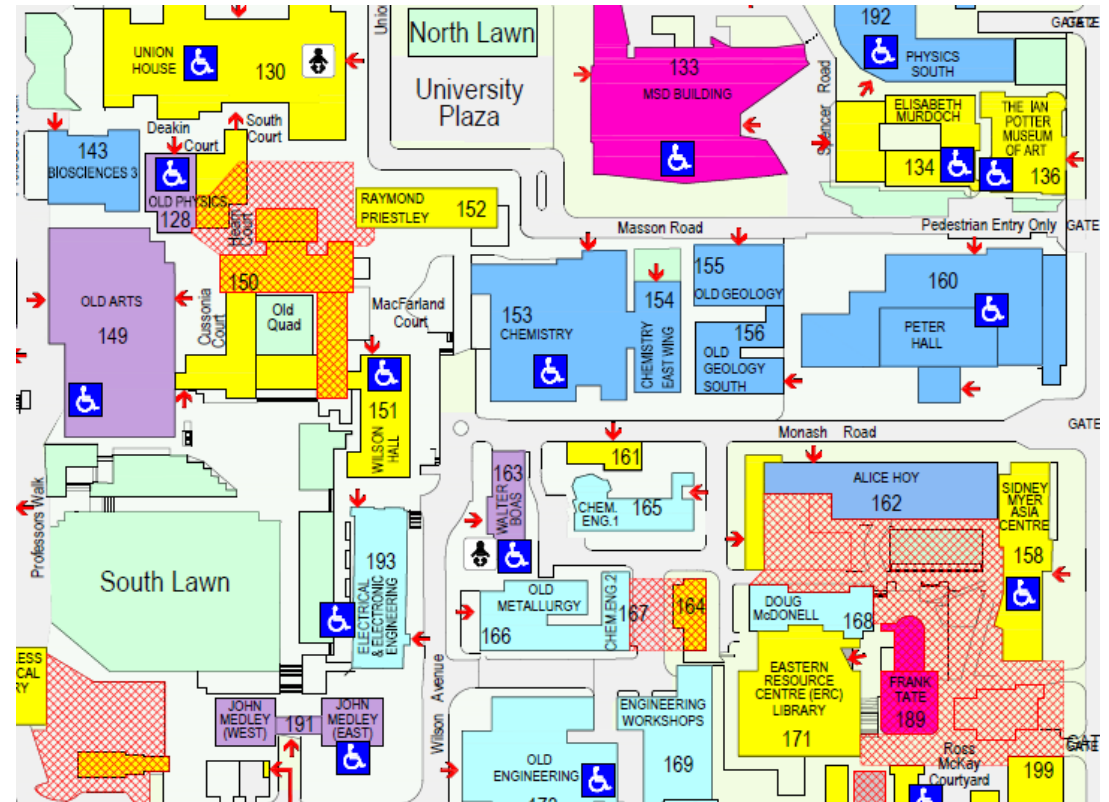
Andres Munoz-Acosta

munoz.m@unimelb.edu.au

Peter Hall Building G.83

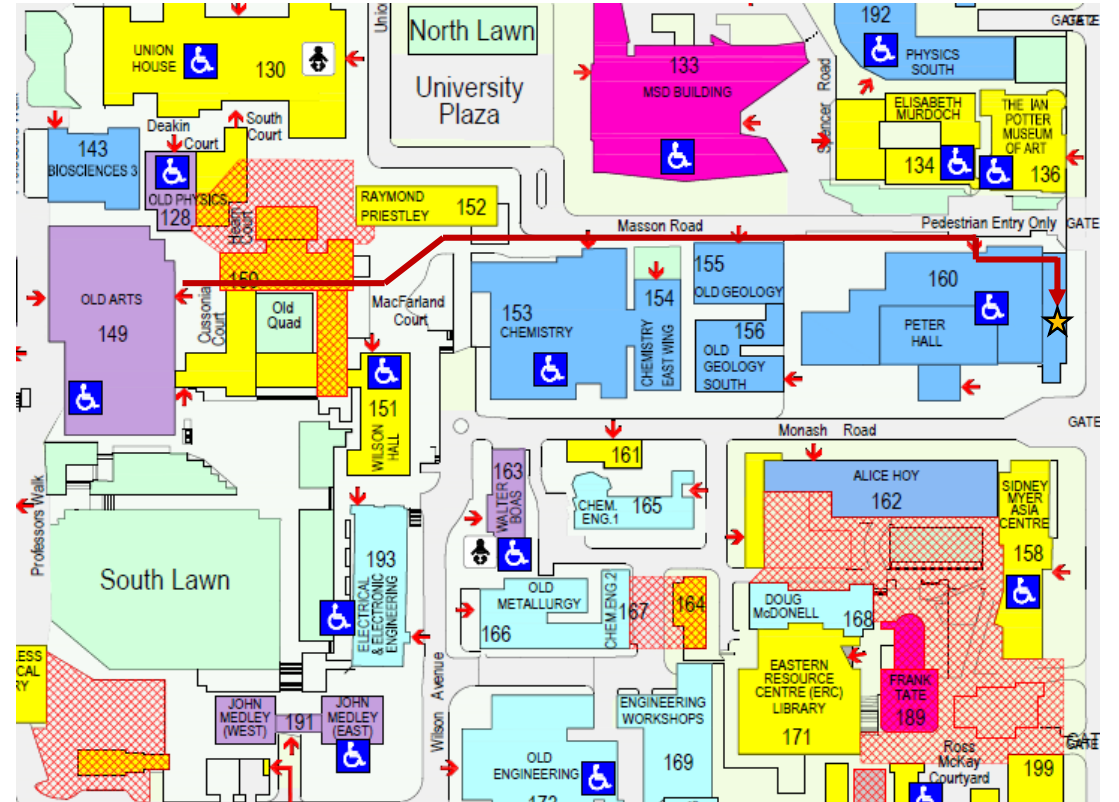
Where to find me?

- My office is at the Peter Hall building (Room G.83)



Where to find me?

- My office is at the Peter Hall building (Room G.83)



Where to find me?

- My office is at the Peter Hall building (Room G.83)
- Consultation hours:
 - Wednesdays 10:00am-11:00am
 - By appointment on Monday/Friday (limited slots)



Heaps and Priority Queues

- The **heap** is a very useful data structure for **priority queues**, used in many algorithms.
- A priority queue is a **set** (or **pool**) of elements.
- An element is injected into the priority queue together with a **priority** (often the key value itself) and elements are ejected according to priority.
- We think of the heap as a **partially ordered binary tree**.
- Since it can easily be maintained as a **complete** tree, the standard implementation uses an array to represent the tree.

The Priority Queue

- As an abstract data type, the priority queue supports the following operations on a “pool” of elements (ordered by some linear order):
 - **find** an item with maximal priority
 - **insert** a new item with associated priority
 - test whether a priority queue is empty
 - **eject** the **largest** element
- Other operations may be relevant, for example:
 - **replace** the maximal item with some new item
 - **construct** a priority queue from a list of items
 - **join** two priority queues

Some Uses of Priority Queues

- **Job scheduling** done by your operating system. The OS will usually have a notion of “importance” of different jobs.
- (Discrete event) **simulation** of complex systems (like traffic, or weather). Here priorities are typically event times.
- Numerical computations involving floating point numbers. Here priorities are measures of computational “error”.
- Many sophisticated algorithms make essential use of priority queues (Huffman encoding and many shortest-path algorithms, for example).

Stacks and Queues as Priority Queues

- Special instances are obtained when we use **time** for priority:
 - If "large" means "late" we obtain the **stack**.
 - If "large" means "early" we obtain the **queue**.

Possible Implementations of the Priority Queue

- Assume priority = key.

Unsorted array or list

Sorted array or list

Heap

INJECT(<i>e</i>)	EJECT()
$O(\log n)$	$O(\log n)$

- How is this accomplished?

The Heap







- A **heap** is a complete binary tree which satisfies the **heap condition**:

Each child has a priority (key) which is no greater than its parent's.

- This guarantees that the root of the tree is a maximal element.
- (Sometimes we talk about this as a **max-heap** – one can equally well have min-heaps, in which each child is no smaller than its parent.)


Heaps and Non-Heaps

- Which of these are heaps?



Heaps as Arrays

- We can utilise the completeness of the tree and place its elements in level-order in an array H .
- Note that the children of node i will be nodes $2i$ and $2i + 1$.




H :

	X	T	O	G	S	M	N	A	E	R	A	I
0	1	2	3	4	5	6	7	8	9	10	11	12

Heaps as Arrays

- This way, the heap condition is very simple:
- For all $i \in \{0, 1, \dots, n\}$, we must have $H[i] \leq H[i/2]$.



H :


	X	T	O	G	S	M	N	A	E	R	A	I
0	1	2	3	4	5	6	7	8	9	10	11	12

Properties of the Heap

- The root of the tree $H[1]$ holds a maximal item; the cost of EJECT is $O(1)$ plus time to restore the heap.
- The height of the heap is $\lfloor \log_2 n \rfloor$.
- Each subtree is also a heap.
- The children of node i are $2i$ and $2i+1$.
- The nodes which happen to be parents are in array positions 1 to $\lfloor n/2 \rfloor$.
- It is easier to understand the heap operations if we think of the heap as a tree.


Injecting a New Item

- Place the new item at the end; then let it "climb up", repeatedly swapping with parents that are smaller:




Injecting a New Item

- Place the new item at the end; then let it "climb up", repeatedly swapping with parents that are smaller:




Injecting a New Item

- Place the new item at the end; then let it "climb up", repeatedly swapping with parents that are smaller:



Building a Heap Bottom-Up


- To construct a heap from an arbitrary set of elements, we can just use the inject operation repeatedly. The construction cost will be $n \log n$. But there is a better way:



- Start with the last parent and move backwards, in level-order. For each parent node, if the largest child is larger than the parent, swap it with the parent.

Building a Heap Bottom-Up


- To construct a heap from an arbitrary set of elements, we can just use the inject operation repeatedly. The construction cost will be $n \log n$. But there is a better way:



- Start with the last parent and move backwards, in level-order. For each parent node, if the largest child is larger than the parent, swap it with the parent.

Building a Heap Bottom-Up


- To construct a heap from an arbitrary set of elements, we can just use the inject operation repeatedly. The construction cost will be $n \log n$. But there is a better way:



- Start with the last parent and move backwards, in level-order. For each parent node, if the largest child is larger than the parent, swap it with the parent.


Building a Heap Bottom-Up: Sifting Down

- Whenever a parent is found to be out of order, let it "sift down" until both children are smaller:




Building a Heap Bottom-Up: Sifting Down

- Whenever a parent is found to be out of order, let it "sift down" until both children are smaller:



Building a Heap Bottom-Up: Sifting Down

- Whenever a parent is found to be out of order, let it "sift down" until both children are smaller:



Turning $H[1] \dots H[n]$ into a Heap, Bottom-Up

```
for  $i \leftarrow \lfloor n/2 \rfloor$  downto 1 do
   $k \leftarrow i$ 
   $v \leftarrow H[k]$ 
   $heap \leftarrow False$ 
  while not  $heap$  and  $2 \times k \leq n$  do
     $j \leftarrow 2 \times k$                                 ▷  $j$  is  $k$ 's left child
    if  $j < n$  then
      if  $H[j] < H[j + 1]$  then
         $j \leftarrow j + 1$                                 ▷  $j$  is  $k$ 's largest child
    if  $v \geq H[j]$  then
       $heap \leftarrow True$ 
    else
       $H[k] \leftarrow H[j]$                                 ▷ Promote  $H[j]$ 
       $k \leftarrow j$ 
   $H[k] \leftarrow v$ 
```


Analysis of Bottom-Up Heap Creation


- For simplicity, assume the heap is a full binary tree: $n = 2^{h+1} - 1$. Here is an upper bound on the number of "down-sifts" needed (consider the root to be at level h , so leaves are at level 0):

$$\sum_{i=1}^h \sum_{\text{nodes at level } i} i = \sum_{i=1}^h i \cdot 2^{h-i} = 2^{h+1} - h - 2$$

- The last equation is easily proved by mathematical induction.
- Note that $2^{h+1} - h - 2 < n$, so we perform at most a linear number of down-sift operations. Each down-sift is preceded by two key comparisons, so the number of comparisons is also linear.
- Hence we have a **linear-time** algorithm for heap creation.


Ejecting a Maximal Element from a Heap

- Here the idea is to swap the root with the last item z in the heap, and then let z "sift down" to its proper place.
- After this, the last element (here shown in green) is no longer considered part of the heap, that is, n is decremented.
- Clearly ejection is $O(\log n)$.




Ejecting a Maximal Element from a Heap

- Here the idea is to swap the root with the last item z in the heap, and then let z "sift down" to its proper place.
- After this, the last element (here shown in green) is no longer considered part of the heap, that is, n is decremented.
- Clearly ejection is $O(\log n)$.




Ejecting a Maximal Element from a Heap

- Here the idea is to swap the root with the last item z in the heap, and then let z "sift down" to its proper place.
- After this, the last element (here shown in green) is no longer considered part of the heap, that is, n is decremented.
- Clearly ejection is $O(\log n)$.



Ejecting a Maximal Element from a Heap

- Here the idea is to swap the root with the last item z in the heap, and then let z "sift down" to its proper place.
- After this, the last element (here shown in green) is no longer considered part of the heap, that is, n is decremented.
- Clearly ejection is $O(\log n)$.



Exercise: Build and Then Deplete a Heap

- First build a heap from the items S, O, R, T, I, N, G.
- Then repeatedly eject the largest, placing it at the end of the heap.

Exercise: Build and Then Deplete a Heap

- First build a heap from the items S, O, R, T, I, N, G.
- Then repeatedly eject the largest, placing it at the end of the heap.
- Anything interesting to notice about the tree that used to hold a heap?

Heapsort

- Heapsort is a $\Theta(n \log n)$ sorting algorithm, based on the idea from this exercise.
- Given an unsorted array $H[1] \dots H[n]$:
- **Step 1:** Turn H into a heap.
- **Step 2:** Apply the eject operation $n-1$ times.

Heapsort

Stage 1 (heap construction)

2 9 **7** 6 5 8

Stage 2 (maximum deletions)

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	5	7

Stage 2 (maximum deletions)

Heapsort

Stage 1 (heap construction)

2 9 **7** 6 5 8
2 **9** 8 6 5 7
2 9 8 6 5 7

Stage 2 (maximum deletions)

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	<u>8</u>	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7

Stage 2 (maximum deletions)

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	<u>8</u>	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7
9	6	8	2	5	7

Stage 2 (maximum deletions)

Heapsort

Stage 1 (heap construction)

2 9 **7** 6 5 8
2 **9** 8 6 5 7
2 9 8 6 5 7
9 **2** 8 6 5 7
9 6 8 2 5 7

Stage 2 (maximum deletions)

9 6 8 2 5 7

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	<u>8</u>	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7
9	6	8	2	5	7

Stage 2 (maximum deletions)

9	6	8	2	5	<u>7</u>
7	6	8	2	5	9

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	8	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7
9	6	8	2	5	7

Stage 2 (maximum deletions)

9	6	8	2	5	<u>7</u>
7	6	8	2	5	9
8	6	7	2	<u>5</u>	9

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	8	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7
9	6	8	2	5	7

Stage 2 (maximum deletions)

9	6	8	2	5	<u>7</u>
7	6	8	2	5	9
8	6	7	2	<u>5</u>	9
5	6	7	2	8	9

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	8	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7
9	6	8	2	5	7

Stage 2 (maximum deletions)

9	6	8	2	5	<u>7</u>
7	6	8	2	5	9
8	6	7	2	<u>5</u>	9
5	6	7	2	8	9
7	6	5	<u>2</u>	8	9
2	6	5	7	8	9

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	8	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7
9	6	8	2	5	7

Stage 2 (maximum deletions)

9	6	8	2	5	<u>7</u>
7	6	8	2	5	9
8	6	7	2	<u>5</u>	9
5	6	7	2	8	9
7	6	5	<u>2</u>	8	9
2	6	5	7	8	9
6	2	<u>5</u>	7	8	9
5	2	6	7	8	9

Heapsort

Stage 1 (heap construction)

2	9	7	6	5	<u>8</u>
2	9	8	<u>6</u>	<u>5</u>	7
2	<u>9</u>	8	6	5	7
9	2	8	<u>6</u>	<u>5</u>	7
9	6	8	2	5	7

Stage 2 (maximum deletions)

9	6	8	2	5	<u>7</u>
7	6	8	2	5	9
8	6	7	2	<u>5</u>	9
5	6	7	2	8	9
7	6	5	<u>2</u>	8	9
2	6	5	7	8	9
6	2	<u>5</u>	7	8	9
5	2	6	7	8	9
5	<u>2</u>	6	7	8	9
2	5	6	7	8	9
2	5	6	7	8	9

Properties of Heapsort

- On average slower than quicksort, but stronger performance guarantee.
- Truly in-place.
- Not stable.

Next lecture

- Transform-and-Conquer
 - Pre-sorting (Levitin Section 6.1)