

COMP90038

Algorithms and Complexity

Lecture 15: Balanced Trees

(with thanks to Harald Søndergaard & Michael Kirley)

Andres Munoz-Acosta

munoz.m@unimelb.edu.au



Peter Hall Building G.83

Recap

- Last week we talked about:
 - Two representations: Heaps and Binary Search Trees
 - An algorithm: Heapsort
 - An strategy: Transform-and-conquer through pre-sorting

Differences between heaps and BSTs

- We have the array [2 3 5 6 7 8 9]
- As a heap:
 - **Each child has a priority (key) which is no greater than its parent's.** This guarantees that the root of the tree is a maximal element.
 - It must be a complete tree (filled top to bottom, left to right)
 - There are many valid heaps!!!
- As a BST:
 - Let the root be r ; then each element in the **left subtree is smaller** than r and each element in the **right sub-tree is larger** than r .
 - A BST is never a heap!!!



Heapsort and Pre-sorting

- Heapsort:
 - Uses the fact that the root of a heap is always the maximal element.
 - It iterates the sequence: Build the heap – eject the root – build the heap – eject the root ...
- Pre-sorting
 - Simplify the problem (through sorting the data) such that an efficient algorithm can be used.

Finding anagrams using pre-sorting

- You are given a very long list of words:

{health, revolution, foolish, garner, drive, praise, traverse, anger, ranger,
... scoop, fall, praise}

- Find the anagrams in the list.
- An approach is to sort each word, sort the list of words, and then find the repeats...

Exercise: Finding Anagrams

health	aehhlt	aerstv	1
revolution	eilnoortvu	aegnr	1
foolish	fhiloos	aegnrr	1
garner	aegnrr	aegnrr	2 (This element is an anagram)
drive	deirv	aehhlt	1
praise	aeiprs	aeiprs	1
traverse	aerstv	afl	1
anger	aegnr	coops	1
ranger	aegnrr	deirv	1
...
scoop	coops	eilnoortvu	1
fall	afl	fhiloos	1
truly	lrtuy	lrtuy	1


Sort each word

Sort the list

Find repeats

Approaches to Balanced Binary Search Trees

- If a BST is "reasonably" balanced, search involves $\Theta(\log n)$ comparisons in the worst case.
- If the BST is "unbalanced", search could be linear.
- To optimise performance, it is important to keep trees "reasonably" balanced.



Approaches to Balanced Binary Search Trees





- Instance simplification approaches: Self-balancing trees
 - **AVL trees**
 - Red-black trees
 - Splay trees
- Representational changes:
 - **2–3 trees**
 - 2–3–4 trees
 - B-trees

AVL Trees

- Named after Adelson-Velsky and Landis.
- Recall that we defined the height of the empty tree as -1.
- For a binary (sub-) tree, let the **balance factor** be the difference between the height of its left sub-tree and that of its right sub-tree.
- An **AVL tree** is a BST in which the balance factor is -1, 0, or 1, for every sub-tree.

AVL Trees: Examples and Counter-Examples


- Which of these are AVL trees?




Building an AVL Tree

- As with standard BSTs, insertion of a new node always takes place at the fringe of the tree.
- If insertion of the new node makes the AVL tree unbalanced (some nodes get balance factors of 2 or -2), transform the tree to regain its balance.
- Regaining balance can be achieved with one or two simple, local transformations, so-called **rotations**.


AVL Trees: R-Rotation




AVL Trees: R-Rotation




AVL Trees: L-Rotation




AVL Trees: L-Rotation




AVL Trees: LR-Rotation




AVL Trees: LR-Rotation




AVL Trees: LR-Rotation




AVL Trees: RL-Rotation



AVL Trees: RL-Rotation




AVL Trees: RL-Rotation



AVL Trees: Where to Perform the Rotation


- Along an unbalanced path, we may have several nodes with balance factor 2 (or -2):



- It is always the **lowest** unbalanced subtree that is re-balanced.


AVL Trees: Where to Perform the Rotation

- Along an unbalanced path, we may have several nodes with balance factor 2 (or -2):




- It is always the **lowest** unbalanced subtree that is re-balanced.

AVL Trees: The Single Rotation, Generally



- This shows an **R-rotation**; an **L-rotation** is similar.

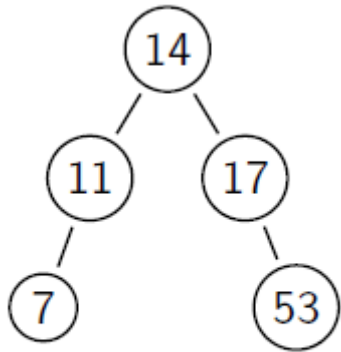
AVL Trees: The Double Rotation, Generally



- This shows an **LR-rotation**; an **RL-rotation** is similar.




Example

- On the tree below, insert the elements {4, 13, 12}







Example

- On the tree below, insert the elements {4, 13, 12}



Example

- On the tree below, insert the elements {4, 13, 12}




Properties of AVL Trees

- The notion of “balance” that is implied by the AVL condition is sufficient to guarantee that the depth of an AVL tree with n nodes is $\Theta(\log n)$.
- For random data, the depth is very close to $\log_2 n$, the optimum.
- In the worst case, search will need at most 45% more comparisons than with a perfectly balanced BST.
- **Deletion** is harder to implement than insertion, but also $\Theta(\log n)$.

Other Kinds of Balanced Trees

- A **red-black tree** is a BSTs with a slightly different concept of “balanced”. Its nodes are coloured red or black, so that
 - No red node has a red child.
 - Every path from the root to the fringe has the same number of black nodes.
- A **splay tree** is a BST which is not only self-adjusting, but also **adaptive**. Frequently accessed items are brought closer to the root, so their access becomes cheaper.




A worst-case red-black tree (the longest path is twice as long as the shortest path).

2–3 Trees

- 2–3 trees and 2–3–4 trees are search trees that allow more than one item to be stored in a tree node.
- A node that holds a single item has two children (or none, if it is a leaf).
- A node that holds two items (a so-called **3-node**) has three children (or none, if it is a leaf).
- And for 2–3–4 trees, a node that holds three items (a **4-node**) has four children (or none, if it is a leaf).
- This allows for a simple way of keeping search trees **perfectly** balanced.


2-Nodes and 3-Nodes



Insertion in a 2–3 Tree

- To insert a key k , pretend that we are searching for k .
- This will take us to a leaf node in the tree, where k should now be inserted; if the node we find there is a 2-node, k can be inserted without further ado.
- Otherwise we had a 3-node, and the two inhabitants, together with k , momentarily form a node with three elements; in sorted order, call them k_1 , k_2 , and k_3 .
- We now **split** the node, so that k_1 and k_3 form their own individual 2-nodes. The middle key, k_2 is **promoted** to the parent node.
- The promotion may cause the parent node to overflow, in which case **it** gets split the same way. The only time the tree's height changes is when the root overflows.

Example: Build a 2–3 Tree from
 $\{9, 5, 8, 3, 2, 4, 7\}$



Exercise: 2–3 Tree Construction

- Build the 2–3 tree that results from inserting these keys, in the given order, into an initially empty tree:

C, O, M, P, U, T, I, N, G

2–3 Tree Analysis

- Worst case search time results when all nodes are 2-nodes. The relation between the number n of nodes and the height h is:

$$n = 1 + 2 + 4 + \dots + 2^h = 2^{h+1} - 1$$

- That is, $\log_2(n+1) = h+1$.

- In the best case, all nodes are 3-nodes:

$$n = 2 + 2 \times 3 + 2 \times 3^2 + \dots + 2 \times 3^h = 3^{h+1} - 1$$

- That is, $\log_3(n+1) = h+1$.

- Hence we have $\log_3(n+1) - 1 \leq h \leq \log_2(n+1) - 1$.

- Useful formula: $\sum_{i=0}^n a^i = \frac{a^{n+1} - 1}{a - 1}$ for $a \neq 1$

Next lecture

- How to buy time, by spending a bit of space.